

Zircon 笔记

许中兴

April 27, 2018

1 Build

让 make 输出详细的编译过程。

```
make V=1 arm64 > make.log
```

在 qemu 里运行 Zircon 的命令为：

```
qemu-system-aarch64 -m 2048 -nographic -net none -smp 4 \
  -kernel /Users/xzx/zircon/build-arm64/qemu-zircon.bin \
  -machine virtualization=true -cpu cortex-a53 \
  -machine virt,gic_version=3 \
  -initrd /Users/xzx/zircon/build-arm64/bootdata.bin \
  -append 'TERM=xterm-256color \
  kernel.entropy-mixin=a7324cb430cbac8... \
  kernel.halt-on-panic=true '
```

可以看出内核文件是 qemu-zircon.bin。我们在 make.log 里找一下这个文件是怎么来的。

```
./build-arm64/tools/mkbootfs -o build-arm64/qemu-zircon.bin \
  ./build-arm64/zircon.bin --header ./build-arm64/qemu-boot-shim.bin \
  --header-align 65536
```

mkbootfs 会把 qemu-zircon.bin 所代表的内核映像写在 offset 为 65536 的地方。在之前写入的是 boot-shim 的内容。qemu-zircon.bin 的布局实际上是由 image.S 决定的。

qemu-boot-shim.bin 的来历：

```
./prebuilt/downloads/gcc/bin/aarch64-elf-objcopy -O binary \
  build-arm64/boot-shim/qemu/boot-shim.elf build-arm64/qemu-boot-shim.bin
```

boot-shim.elf 的来历：

```
./prebuilt/downloads/gcc/bin/aarch64-elf-ld -nostdlib --build-id -z noexecstack \
  -z max-page-size=4096 --gc-sections --emit-relocs --build-id=none \
  ./build-arm64/boot-shim/qemu/boot-shim.S.o \
  ./build-arm64/boot-shim/qemu/boot-shim.c.o \
```

```

./build-arm64/boot-shim/qemu/debug.c.o \
./build-arm64/boot-shim/qemu/devicetree.c.o \
./build-arm64/boot-shim/qemu/util.c.o \
-T kernel/target/arm64/boot-shim/boot-shim.lds \
-o build-arm64/boot-shim/qemu/boot-shim.elf

```

boot-shim.S 中 `_start` 的地址是 0。所以 qemu-zircon.bin 加载之后，自然就从 boot-shim.S 的 `_start` 开始执行。那么，内核的 `_start` 又是如何找到的呢？注意到有这样一段代码：

```

// x0: pointer to device tree
// x1: pointer to kernel bootdata container
bl      boot_shim

// kernel entry point is returned in x0
mov     tmp, x0

// pass bootdata to kernel in x0
adr     x0, bootdata_return
ldr     x0, [x0]

br      tmp

```

也就是说，`boot_shim` 函数应该会找到 kernel 的 entry point。我们去看一下。

```
return kernel_base + kernel->data_kernel.entry64;
```

这里 `kernel_base` 就是 `KERNEL_ALIGN=65536`。`boot_shim` 函数从 bootdata 中获得入口地址。bootdata 是 `mkbootfs` 这个工具填写到启动映像里的。在 `image.S` 的头上就是 bootdata。

```

// BOOTDATA_KERNEL payload (bootdata_kernel_t)
DATA(_bootdata_kernel_payload)
    // The boot-shim code expects this to be an offset from the beginning
    // of the load image, whatever the kernel's virtual address.
    .quad IMAGE_ELF_ENTRY - _bootdata_file_header
    .quad 0
END_DATA(_bootdata_kernel_payload)

```

从而，上面的 `br tmp` 能够正确跳转到内核入口。
zircon.bin 的来历：

```

./prebuilt/downloads/gcc/bin/aarch64-elf-objcopy \
-O binary build-arm64/zircon-image.elf build-arm64/zircon.bin

```

zircon-image.elf 的来历：

```
./prebuilt/downloads/gcc/bin/aarch64-elf-ld -nostdlib --build-id -z noexecstack \
    -z max-page-size=4096 --gc-sections --build-id=none \
-o build-arm64/zircon-image.elf -T kernel/image.ld \
    --just-symbols ./build-arm64/zircon.elf \
./build-arm64/kernel-vars.ld ./build-arm64/zircon.image.o
```

链接脚本 image.ld 内容:

```
ENTRY (IMAGE_ELF_ENTRY)

SECTIONS {
    . = IMAGE_LOAD_START;

    .load_image : {
        KEEP(*(.text))
    } :load_image

    /*
     * When a boot loader is actually using the ELF headers, it needs to
     * know how much memory to reserve after the load image (p_filesz is
     * the load image, and p_memsz > p_filesz to indicate the extra space
     * to reserve). This ensures that the segment has the right p_memsz.
     */
    .bss : {
        . += ABSOLUTE (IMAGE_MEMORY_END) - ABSOLUTE (.);
        ASSERT (ABSOLUTE (.) == ABSOLUTE (IMAGE_MEMORY_END), "image layout bad");
    }
}

PHDRS {
    load_image PT_LOAD FLAGS(7); /* PF_R|PF_W|PF_X */
}
```

ENTRY 指定了程序入口的虚拟地址，它对应于 ELF Header 中的 entry 项。
IMAGE_ELF_ENTRY 的定义在 kernel/arch/arm64/start.S 中。

```
// This symbol is used by image.S.
.global IMAGE_ELF_ENTRY
IMAGE_ELF_ENTRY = _start
```

IMAGE_LOAD_START 是在 kernel.ld 中定义的符号，它的值是-4GB

```
. = KERNEL_BASE;
PROVIDE_HIDDEN (__code_start = .);

...

IMAGE_LOAD_START = __code_start;
```

KERNEL_BASE 定义在 kernel/arch/arm64/rules.mk

```
KERNEL_BASE := 0xffffffff00000000
```

注意：整个内核映像的起始虚拟地址是-4GB，但是入口并不是这个地址，而是 start.S 文件中的 `_start`。bootloader 在将内核加载到内存中之后，应该设法跳转到 `_start` 处。我们到后面看看是如何跳转的。也许在 mkbootfs 中有线索。这里的疑问是，objcopy -O binary 实质上只是把 .text 拷贝出来，丢弃了 entry 信息，因为 entry 信息是保存在 elf header 里面的。

(暂时搞不清楚 qemu 是如何跳转到 `_start` 的。boot-shim 里也有 `_start` 符号。)

这里 zircon.elf 只是用来获得 symbol 的地址，并不添加到 zircon-image.elf 里。真正的内核映像文件是 zircon.image.o

```
./prebuilt/downloads/gcc/bin/aarch64-elf-gcc \
-O2 -g -fdebug-prefix-map=/Users/xzx/zircon=. -finline \
  -include ./build-arm64/config-global.h -Wall -Wextra -Wno-multichar \
  -Werror -Wno-error=deprecated-declarations -Wno-unused-parameter \
  -Wno-unused-function -Werror=unused-label -Werror=return-type -fno-common \
  -ffunction-sections -fdata-sections -Wno-nonnull-compare -ffreestanding \
  -include ./build-arm64/config-kernel.h -Wformat=2 -Wvla -Wformat-signedness \
  -fno-exceptions -fno-unwind-tables -fno-omit-frame-pointer -mgeneral-regs-only \
  -fPIE -include kernel/include/hidden.h -mcpu=cortex-a53 -ffixed-x18 \
  -Isystem/public -Isystem/private -I./build-arm64/gen/global/include \
  -I./build-arm64 -Ikernel/include -Ikernel/platform/generic-arm/include \
  -Ikernel/arch/arm64/include -Ikernel/top/include \
  -Ikernel/dev/hdcp/amlogic_s912/include \
  -Ikernel/dev/interrupt/arm_gic/common/include \
  -Ikernel/dev/interrupt/arm_gic/v2/include \
  -Ikernel/dev/interrupt/arm_gic/v2/include \
  -Ikernel/dev/interrupt/arm_gic/v3/include \
  -Ikernel/dev/iommu/dummy/include -Ikernel/dev/pcie/include \
  -Ikernel/dev/pdev/include -Ikernel/dev/pdev/power/include \
  -Ikernel/dev/power/hisi/include -Ikernel/dev/psci/include \
  -Ikernel/dev/timer/arm_generic/include -Ikernel/dev/uart/amlogic_s905/include \
  -Ikernel/dev/uart/nxp-imx/include -Ikernel/dev/uart/pl011/include \
  -Ikernel/kernel/include -Isystem/ulib/bitmap/include \
  -Ikernel/lib/bitmap/include -Ikernel/lib/cbuf/include \
  -Ikernel/lib/debugcommands/include -Ikernel/lib/debuglog/include \
  -Ikernel/lib/ktrace/include -Ikernel/lib/memory_limit/include \
  -Ikernel/lib/mtrace/include -Ikernel/lib/userboot/include \
  -Ikernel/lib/version/include -Ikernel/object/include \
  -Ikernel/platform/include -Ikernel/syscalls/include \
  -Ikernel/target/include -Ikernel/tests/include \
  -Ikernel/dev/interrupt/include -Ikernel/dev/pdev/interrupt/include \
  -Ikernel/dev/pdev/uart/include -Ikernel/dev/udisplay/include \
```

```

-Ikernel/lib/console/include -Ikernel/lib/counters/include \
-Ikernel/lib/crypto/include -Ikernel/lib/debug/include \
-Isystem/ulib/explicit-memory/include -Ikernel/lib/explicit-memory/include \
-Ikernel/lib/fbl/include -Isystem/ulib/fbl/include \
-Ikernel/lib/fbl/include -Ikernel/lib/fixed_point/include \
-Ikernel/lib/header_tests/include -Ikernel/lib/heap/include \
-Ikernel/lib/heap/include -Ikernel/lib/hypervisor/include \
-Ikernel/lib/libc/include -Ikernel/lib/oom/include -Ikernel/lib/pci/include \
-Ikernel/lib/pci/include -Ikernel/lib/pow2_range_allocator/include \
-Isystem/ulib/region-alloc/include -Ikernel/lib/region-alloc/include \
-Ikernel/lib/unittest/include -Ikernel/lib/user_copy/include \
-Ikernel/lib/vdso/include -Isystem/ulib/zxcpp/include \
-Ikernel/lib/zxcpp/include -Ikernel/vm/include \
-Ikernel/arch/arm64/hypervisor/include -Ikernel/dev/hw_rng/include \
-Ikernel/lib/gfx/include -Ikernel/lib/gfxconsole/include \
-Ikernel/lib/heap/cmpctmalloc/include -Ikernel/lib/io/include \
-Isystem/ulib/pretty/include -Ikernel/lib/pretty/include \
-Ithird_party/ulib/cryptolib/include -Ithird_party/lib/cryptolib/include \
-Ithird_party/lib/jitterentropy/include -Ithird_party/lib/jitterentropy/include \
-Ithird_party/ulib/qrcodegen/include -Ithird_party/lib/qrcodegen/include \
-Ithird_party/ulib/uboringssl/include -Ithird_party/lib/uboringssl/include \
-I./build-arm64 \
-c kernel/arch/arm64/image.S -MD -MP -MT build-arm64/zircon.image.o \
-MF build-arm64/zircon.image.d -o build-arm64/zircon.image.o

```

image.S 文件里嵌入了真正的内核 zircon.elf.bin

```

./prebuilt/downloads/gcc/bin/aarch64-elf-objcopy \
-O binary build-arm64/zircon.elf build-arm64/zircon.elf.bin

```

真正的 kernel 是 zircon.elf

```

./prebuilt/downloads/gcc/bin/aarch64-elf-ld -nostdlib --build-id -z noexecstack \
-z max-page-size=4096 --gc-sections --emit-relocs -T kernel/kernel.ld \
build-arm64/kernel-vars.ld \
build-arm64/kernel/platform/generic-arm/generic-arm.mod.o \
build-arm64/kernel/arch/arm64/arm64.mod.o build-arm64/kernel/top/top.mod.o \
build-arm64/kernel/dev/hdcp/amlogic_s912/amlogic_s912.mod.o \
build-arm64/kernel/dev/interrupt/arm_gic/common/common.mod.o \
build-arm64/kernel/dev/interrupt/arm_gic/v2/v2.mod.o \
build-arm64/kernel/dev/interrupt/arm_gic/v3/v3.mod.o \
build-arm64/kernel/dev/iommu/dummy/dummy.mod.o \
build-arm64/kernel/dev/pcie/pcie.mod.o \
build-arm64/kernel/dev/pdev/pdev.mod.o \
build-arm64/kernel/dev/pdev/power/power.mod.o \
build-arm64/kernel/dev/power/hisi/hisi.mod.o \
build-arm64/kernel/dev/psci/psci.mod.o \

```

```

build-arm64/kernel/dev/timer/arm_generic/arm_generic.mod.o \
build-arm64/kernel/dev/uart/amlogic_s905/amlogic_s905.mod.o \
build-arm64/kernel/dev/uart/nxp-imx/nxp-imx.mod.o \
build-arm64/kernel/dev/uart/pl011/pl011.mod.o \
build-arm64/kernel/kernel/kernel.mod.o \
build-arm64/kernel/lib/bitmap/bitmap.mod.o \
build-arm64/kernel/lib/cbuf/cbuf.mod.o \
build-arm64/kernel/lib/debugcommands/debugcommands.mod.o \
build-arm64/kernel/lib/debuglog/debuglog.mod.o \
build-arm64/kernel/lib/ktrace/ktrace.mod.o \
build-arm64/kernel/lib/memory_limit/memory_limit.mod.o \
build-arm64/kernel/lib/mtrace/mtrace.mod.o \
build-arm64/kernel/lib/userboot/userboot.mod.o \
build-arm64/kernel/lib/version/version.mod.o \
build-arm64/kernel/object/object.mod.o \
build-arm64/kernel/platform/platform.mod.o \
build-arm64/kernel/syscalls/syscalls.mod.o \
build-arm64/kernel/target/target.mod.o \
build-arm64/kernel/tests/tests.mod.o \
build-arm64/kernel/dev/interrupt/interrupt.mod.o \
build-arm64/kernel/dev/pdev/interrupt/interrupt.mod.o \
build-arm64/kernel/dev/pdev/uart/uart.mod.o \
build-arm64/kernel/dev/udisplay/udisplay.mod.o \
build-arm64/kernel/lib/console/console.mod.o \
build-arm64/kernel/lib/counters/counters.mod.o \
build-arm64/kernel/lib/crypto/crypto.mod.o \
build-arm64/kernel/lib/debug/debug.mod.o \
build-arm64/kernel/lib/explicit-memory/explicit-memory.mod.o \
build-arm64/kernel/lib/fbl/fbl.mod.o \
build-arm64/kernel/lib/fixed_point/fixed_point.mod.o \
build-arm64/kernel/lib/header_tests/header_tests.mod.o \
build-arm64/kernel/lib/heap/heap.mod.o \
build-arm64/kernel/lib/hypervisor/hypervisor.mod.o \
build-arm64/kernel/lib/libc/libc.mod.o \
build-arm64/kernel/lib/oom/oom.mod.o \
build-arm64/kernel/lib/pci/pci.mod.o \
build-arm64/kernel/lib/pow2_range_allocator/pow2_range_allocator.mod.o \
build-arm64/kernel/lib/region-alloc/region-alloc.mod.o \
build-arm64/kernel/lib/unittest/unittest.mod.o \
build-arm64/kernel/lib/user_copy/user_copy.mod.o \
build-arm64/kernel/lib/vdso/vdso.mod.o \
build-arm64/kernel/lib/zxcpp/zxcpp.mod.o \
build-arm64/kernel/vm/vm.mod.o \
build-arm64/kernel/arch/arm64/hypervisor/hypervisor.mod.o \
build-arm64/kernel/dev/hw_rng/hw_rng.mod.o \
build-arm64/kernel/lib/gfx/gfx.mod.o \

```

```

build-arm64/kernel/lib/gfxconsole/gfxconsole.mod.o \
build-arm64/kernel/lib/heap/cmpctmalloc/cmpctmalloc.mod.o \
build-arm64/kernel/lib/io/io.mod.o \
build-arm64/kernel/lib/pretty/pretty.mod.o \
build-arm64/third_party/lib/cryptolib/cryptolib.mod.o \
build-arm64/third_party/lib/jitterentropy/jitterentropy.mod.o \
build-arm64/third_party/lib/qrcodegen/qrcodegen.mod.o \
build-arm64/third_party/lib/uboringssl/uboringssl.mod.o \
-o build-arm64/zircon.elf

```

我们先看 kernel.ld:

```

/*
 * Symbols used in the kernel proper are defined with PROVIDE_HIDDEN:
 * HIDDEN because everything in the kernel is STV_HIDDEN to make it
 * clear that direct PC-relative references should be generated in PIC;
 * PROVIDE because their only purpose is to satisfy kernel references.
 */

SECTIONS
{
    . = KERNEL_BASE;
    PROVIDE_HIDDEN(__code_start = .);

    /*
     * This just leaves space in the memory image for the boot headers.
     * The actual boot header will be constructed in image.S, which see.
     */
    .text.boot0 : {
        /*
         * Put some data in, or else the linker makes it a SHT_NOBITS
         * section and that makes objcopy -O binary skip it in the image.
         */
        LONG(0xdeadbeef);
        . += BOOT_HEADER_SIZE - 4;
    } :code

    . = ALIGN(8);
    .buildsig : {
        PROVIDE_HIDDEN(buildsig = .);
        BYTE(0x42); BYTE(0x53); BYTE(0x49); BYTE(0x47); /* BSIG */
        BYTE(0x53); BYTE(0x54); BYTE(0x52); BYTE(0x54); /* STRT */
        /*
         * The self-pointer gives a local basis to compute the relative
         * positions of the lk_version_t and .note.gnu.build-id pointers
         * without already knowing the kernel's virtual address base.
         */
    }
}

```

```

        QUAD(buildsig);
        QUAD(version);
        QUAD(__build_id_note_start);
        BYTE(0x42); BYTE(0x53); BYTE(0x49); BYTE(0x47); /* BSIG */
        BYTE(0x45); BYTE(0x4e); BYTE(0x44); BYTE(0x53); /* ENDS */
    }

/*
 * This is separate from .text just so gen-kaslr-relocs.sh can match
 * it. The relocation processor skips this section because this code
 * all runs before the boot-time fixups are applied and has its own
 * special relationship with the memory layouts.
 */
.text.boot : {
    *(.text.boot)
}

/*
 * This is separate from .text just so gen-kaslr-relocs.sh can match
 * it. This section contains movabs instructions that get 64-bit
 * address fixups in place. This is safe because this code is never
 * used until long after fixups have been applied. In general, the
 * script will refuse to handle fixups in text (i.e. code) sections.
 */
.text.bootstrap16 : {
    *(.text.bootstrap16)
}

.text : {
    *(.text* .sram.text)
    *(.gnu.linkonce.t.*)
}

PROVIDE_HIDDEN(__code_end = .);

. = ALIGN(4096);
PROVIDE_HIDDEN(__rodata_start = .);

/*
 * These are page-aligned, so place them first.
 */
.rodata.rodso_image : {
    *(.rodata.rodso_image.*)
}

.note.gnu.build-id : {

```



```

        PROVIDE_HIDDEN(__build_id_note_start = .);
        *(.note.gnu.build-id)
        PROVIDE_HIDDEN(__build_id_note_end = .);
} :rodata :note

/*
 * The named sections starting with kcountdesc are sorted
 * by name so that tools can provide binary search lookup
 * for k_counter_desc::name[] variables.
 */
.kcounter.desc : ALIGN(8) {
    PROVIDE_HIDDEN(kcountdesc_begin = .);
    KEEP(*(SORT_BY_NAME(kcountdesc.*)))
    PROVIDE_HIDDEN(kcountdesc_end = .);
} :rodata

.rodata : {
    *(.rodata* .gnu.linkonce.r.*)
} :rodata

/*
 * When compiling PIC, the compiler puts things into sections it
 * thinks need to be writable until after dynamic relocation. In
 * the kernel, these things all go into the read-only segment. But
 * to the linker, they are writable and so the default "orphans"
 * placement would put them after .data instead of here. That's bad
 * both because we want these things in the read-only segment (the
 * kernel's self-relocation applies before the read-only-ness starts
 * being enforced anyway), and because the orphans would wind up
 * being after the __data_end symbol (see below).
 *
 * Therefore, we have to list all the special-case sections created
 * by __SECTION("foo") uses in the kernel that are RELRO candidates,
 * i.e. things that have address constants in their initializers.
 * All such uses in the source use sections named ".data.rel.ro.foo"
 * instead of just "foo" specifically to ensure we write them here.
 * This avoids the magic linker behavior for an "orphan" section
 * called "foo" of synthesizing "__start_foo" and "__stop_foo"
 * symbols when the section name has no . characters in it, and so
 * makes sure we'll get undefined symbol references if we omit such
 * a section here. The magic linker behavior is nice, but it only
 * goes for orphans, and we can't abide the default placement of
 * orphans that should be RELRO.
 */
.data.rel.ro : ALIGN(8) {
    PROVIDE_HIDDEN(__start_commands = .);

```

```

KEEP*(.data.rel.ro.commands))
PROVIDE_HIDDEN(__stop_commands = .);

PROVIDE_HIDDEN(__start_ktrace_probe = .);
KEEP*(.data.rel.ro.ktrace_probe))
PROVIDE_HIDDEN(__stop_ktrace_probe = .);

PROVIDE_HIDDEN(__start_lk_init = .);
KEEP*(.data.rel.ro.lk_init))
PROVIDE_HIDDEN(__stop_lk_init = .);

PROVIDE_HIDDEN(__start_lk_pdev_init = .);
KEEP*(.data.rel.ro.lk_pdev_init))
PROVIDE_HIDDEN(__stop_lk_pdev_init = .);

PROVIDE_HIDDEN(__start_unittest_testcases = .);
KEEP*(.data.rel.ro.unittest_testcases))
PROVIDE_HIDDEN(__stop_unittest_testcases = .);

*(.data.rel.ro* .gnu.linkonce.d.rel.ro.*)
}

.init_array : ALIGN(8) {
    PROVIDE_HIDDEN(__init_array_start = .);
    KEEP*(.init_array.ctors))
    PROVIDE_HIDDEN(__init_array_end = .);
}

/*
 * Any read-only data "orphan" sections will be inserted here.
 * Ideally we'd put those into the .rodata output section, but
 * there isn't a way to do that that guarantees all same-named
 * input sections collect together as a contiguous unit, which
 * is what we need them for. Linkers differ in how they'll
 * place another dummy section here relative to the orphans, so
 * there's no good way to define __rodata_end to be exactly the
 * end of all the orphans sections. But the only use we have
 * for __rodata_end is to round it up to page size anyway, so
 * just define it inside the .data section below, which is
 * exactly the end of the orphans rounded up to the next page.
 */

.data : ALIGN(4096) {
    PROVIDE_HIDDEN(__rodata_end = .);
    PROVIDE_HIDDEN(__data_start = .);
    *(.data .data* .gnu.linkonce.d.*)

```

```

/*
 * Make sure the total file size is aligned to 8 bytes so the image
 * can go into a BOOTDATA container, which requires 8-byte alignment.
 */
. = ALIGN(8);
PROVIDE_HIDDEN(__data_end = .);
} :data

/*
 * Any writable orphan sections would be inserted here.
 * But there's no way to put the __data_end symbol after
 * them, so we cannot allow any such cases. There is no
 * good way to assert that, though.
 */

.bss : ALIGN(4096) {
    PROVIDE_HIDDEN(__bss_start = .);

    /*
     * See kernel/include/lib/counters.h; the KCOUNTER macro defines a
     * kcounter.NAME array in the .bss.kcounter.NAME section that
     * allocates SMP_MAX_CPUS counter slots. Here we collect all those
     * together to make up the kcounters_arena contiguous array. There
     * is no particular reason to sort these, but doing so makes them
     * line up in parallel with the sorted .kcounter.desc section.
     */
    . = ALIGN(8);
    PROVIDE_HIDDEN(kcounters_arena = .);
    KEEP(*(SORT_BY_NAME(.bss.kcounter.*)))

    /*
     * Sanity check that the aggregate size of kcounters_arena
     * SMP_MAX_CPUS slots for each counter. The k_counter_desc structs
     * in .kcounter.desc are 8 bytes each, which matches the size of a
     * single counter. (It's only for this sanity check that we need
     * to care how big k_counter_desc is.)
     */
    ASSERT(. - kcounters_arena == sizeof(.kcounter.desc) * SMP_MAX_CPUS,
           "kcounters_arena size mismatch");

    *(.bss*)
    *(.gnu.linkonce.b.*)
    *(COMMON)
}

/*

```

```

    * Any SHT_NOBITS (.bss-like) sections would be inserted here.
    */

    . = ALIGN(4096);
    PROVIDE_HIDDEN(_end = .);
}

PHDRS
{
    code PT_LOAD FLAGS(5); /* PF_R|PF_X */
    rodata PT_LOAD FLAGS(4); /* PF_R */
    data PT_LOAD FLAGS(6); /* PF_R|PF_W */
    note PT_NOTE FLAGS(4); /* PF_R */
}

/*
 * This is not actually used since the entry point is set in image.ld,
 * but it prevents the linker from warning about using a default address
 * and it keeps --gc-sections from removing .text.boot.
 */
ENTRY(IMAGE_ELF_ENTRY)

/*
 * These special symbols below are made public so they are visible via
 * --just-symbols to the link of image.S.
 */

IMAGE_LOAD_START = __code_start;
IMAGE_LOAD_END = __data_end;
IMAGE_MEMORY_END = _end;

```

2 start.S

内核真正的入口是 start.S 里面的 `_start`。
一些基础知识。

http://refspecs.linuxfoundation.org/LSB_3.0.0/LSB-Core-generic/LSB-Core-generic/ehframechpt

The `.eh_frame` section shall contain 1 or more Call Frame Information (CFI) records. The number of records present shall be determined by size of the section as contained in the section header. Each CFI record contains a Common Information Entry (CIE) record followed by 1 or more Frame Description Entry (FDE) records. Both CIEs and FDEs shall be aligned to an addressing unit sized boundary.

`.cfi_startproc` is used at the beginning of each function that should have an entry in `.eh_frame`. It initializes some internal data structures. Don't forget to close the function by `.cfi_endproc`.

Unless `.cfi_startproc` is used along with `parameter simple` it also emits some architecture dependent initial CFI instructions.